

Quiz Submissions - Week 7 Quiz



Timothy Holmes (username: THOLME15)

Attempt 2

Written: Oct 26, 2020 8:22 PM - Oct 26, 2020 8:26 PM

Submission View

Released: Sep 7, 2020 10:15 PM

Network Programming

Question 1

1 / 1 point

What is a first-level domain name?

- A. These are the domain names of the most visited websites on the Internet.
- B. The six original end of domain names: gov, edu, com, mil, org, and net.
- C. Also called Top-Level Domains (TLD), they are domain names living as direct children of the DNS root.
- D. A domain name that contains a dot (.).

☐ Answer A.

☐ Answer B.

☒ Answer C.

☐ Answer D.

[▶ View Feedback](#)

Question 2

1 / 1 point

The main purpose of the Domain Naming System is to map domain names (text addresses) to IP addresses. Where is that database stored?

- A. There's no single place, it is heavily distributed.
- B. On an ISC (Internet Systems Consortium) server with "well-known" IPv4 address 192.5.5.241 (IPv6: 2001:500:2f::f).
- C. On the ISP (Internet Service Provider) DNS server.
- D. On Google's DNS server (8.8.8.8).

✓ ☒ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

Question 3

1 / 1 point

In an iterative DNS query, how many DNS servers can be contacted to resolve `www.cdm.depaul.edu`?

A. 1 B. 2 C. 3 D. 4

☐ Answer A.

☐ Answer B.

☐ Answer C.

✓ ☒ Answer D.

▶ [View Feedback](#)

Question 4

1 / 1 point

To resolve an address, the libc will contact a DNS server. How does the libc, on a Unix-like system such as Linux, know the IP address of the DNS server?

- A. The local router answers with the address of a DNS server when a connection is attempted to the "well-known" address 127.0.0.1.
- B. A list of DNS servers is provided by the file `/etc/resolv.conf`.
- C. A default DNS server is hardcoded in the libc.
- D. The libc does not know of a DNS server address: to resolve an address, it makes a system call and the kernel will use the DNS server it received when it was given an IP address.

☐ Answer A.

✓ ☒ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

Question 5

1 / 1 point

DNS can be seen as a function mapping domain names to sets of IPs (since a domain name can have multiple IPs). Calling f this function, is the following situation possible?

$$f(\text{domain1}) = \{\text{ip1}, \text{ip2}\}$$

$$f(\text{domain2}) = \{\text{ip1}\}$$

$$f(\text{domain3}) = \{\text{ip2}\}$$

- A. No; since ip1 serves domain1, it cannot serve another domain.
- B. No; since both ip1 and ip2 serve domain1, they should also *both* serve domain2 and domain3.
- C. Yes.
- D. Yes, but it is a temporary situation due to outdated DNS caches.

☐ Answer A.

☐ Answer B.

✓ ☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

Question 6

1 / 1 point

HTTP relies on TCP, which relies on IP, which relies on Ethernet (or 802.11 Wi-Fi, for instance). All of these are *protocols*. Where is the notion of *port* defined? (E.g., HTTP works on port 80.)

- A. At the application level, i.e., it is defined by HTTP.
- B. At the transport level, i.e., it is defined by TCP.
- C. At the network level, i.e., it is defined by IP.
- D. At the data link level, i.e., it is defined by Ethernet.

☐ Answer A.

✓ ☐ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

Question 7

1 / 1 point

(TCP) connections are port-to-port. When a client connects to the HTTP port of a server (80), what is the port allocated on the client?

- A. A random, so-called *ephemeral* port.
- B. Port 80.
- C. Port 80 + 1024, since ports below 1024 are reserved.
- D. None; in a client-server connection, only the server needs a port.

✓ ☐ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

Question 8

1 / 1 point

(TCP) connections are said to be process-to-process. When a connection is made to a server using a port number (say 80), how is the correct process notified?

- A. All processes on the server-side are notified, and they should test that the target port is the one they expect.
- B. The TCP packet should contain the process ID of the correct process on the server side.
- C. A port can only be allocated by *one* process, so the server-side kernel knows which process is being targeted.

☐ Answer A.

☐ Answer B.

✓ ☒ Answer C.

▶ [View Feedback](#)

Question 9

1 / 1 point

Applications read from and write to sockets. In C, what *is* a socket?

- A. A pointer to **char**, that acts as a buffer to the connection.
- B. A file descriptor.
- C. An object of type **FILE***.
- D. An object of type **struct sockaddr**.

☐ Answer A.

✓ ☒ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

Question 10

1 / 1 point

What is a `struct sockaddr`?

- A. A socket that can be read from or written to.
- B. A structure with a large number of fields, that covers most common socket addressing uses.
- C. A protocol-agnostic, generic type to address remote sockets.
- D. An IPv4 addressing structure.

☐ Answer A.

☐ Answer B.

✓ ☒ Answer C.

☐ Answer D.

▶ [View Feedback](#)

Question 11

1 / 1 point

What is the succession of calls necessary to open a client socket, and make a connection request?

- A. `getaddrinfo`, `connect`, `socket`.
- B. `getaddrinfo`, `socket`, `connect`.
- C. `getaddrinfo`, `connect`, `bind`, `socket`.
- D. `getaddrinfo`, `socket`, `bind`, `connect`.

☐ Answer A.

✓ ☒ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

Question 12

1 / 1 point

What is the succession of calls necessary to open a listening socket, and wait for connections?

- A. getaddrinfo, socket, bind, listen, accept.
- B. getaddrinfo, bind, socket, accept, listen.
- C. getaddrinfo, socket, bind, accept, listen.
- D. getaddrinfo, socket, listen, bind, accept.

✓ ☒ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

Question 13

1 / 1 point

Assuming that a server only *reads* from its clients, how is the server notified that a client has closed its end of the connection?

- A. It will receive a EOF value (typically -1) on the next read on the socket.
- B. It will receive a signal on the next read on the socket.
- C. It will receive a signal as soon as the client closes.

✓ ☒ Answer A.

☐ Answer B.

☐ Answer C.

[▶ View Feedback](#)**Question 14****1 / 1 point**

Where is DNS resolving implemented on a UNIX-like system, such as Linux?

- A. In a POSIX library known as `libdns`.
- B. In the kernel.
- C. In the network adapter.
- D. In the `libc`.

☐ Answer A.

☐ Answer B.

☐ Answer C.

✓ ☒ Answer D.

[▶ View Feedback](#)**Question 15****1 / 1 point**

How is most of the information retrieved by `getaddrinfo` returned to the caller?

- A. As an automatically allocated linked list of `struct addrinfo`.
- B. Using a global variable `addrinfo` of type `struct addrinfo`.
- C. It is written in a string buffer provided by the caller, of a given maximum size. This buffer will contain a string describing the info.
- D. It does not return information explicitly, but the information is used implicitly in subsequent calls to socket interface functions.

✓ ☒ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

[▶ View Feedback](#)**Question 16****1 / 1 point**

In a normal application relying on DNS, assume there is a call:

```
connect(fd, addr, addr_len);
```

How were the values of `addr` and `addr_len` typically found?

- A. They are computed using various bit masks, and some macros are provided to simplify this.
- B. They are constants found in the header file `sys/socket.h`.
- C. They are the obvious values, for instance `addr` would be the string `"127.0.0.1"`, and in that case, `addr_len` would be 9.
- D. They are returned by `getaddrinfo`, which should be called beforehand.

☐ Answer A.

☐ Answer B.

☐ Answer C.

☒ Answer D.

[▶ View Feedback](#)**Question 17****1 / 1 point**

The function `getaddrinfo` can be seen as converting “human-readable” addressing information to a `struct sockaddr`. What function provides the converse direction?

- A. `getinfoaddr`.
- B. `getnameinfo`.
- C. `getsockaddrinfo`.
- D. `getaddrinfo_r`.

☐ Answer A.

☒ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

Question 18

1 / 1 point

Among the 5 fundamental socket system calls, only one is usually used by both the client and the server. Which one?

A. socket

B. bind

C. listen

D. accept

E. connect

✓ ☒ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

☐ Answer E.

▶ [View Feedback](#)

Question 19

1 / 1 point

The call to `socket(domain, type, protocol)` returns a socket descriptor. How does one indicate that the socket is supposed to be a client type, rather than server type?

A. Using a specific value for domain.

B. Using a specific value for type.

C. Using a specific value for protocol.

D. The kernel does not differentiate between the two types of sockets.

- ☐ Answer A.
- ☐ Answer B.
- ☐ Answer C.
- ✓ ☒ Answer D.

▶ [View Feedback](#)

Question 20

1 / 1 point

What is the role of the syscall `bind`?

- A. Mark a socket descriptor as *passive*, that is, mark it as a socket that will be used for incoming connections.
- B. It creates a socket descriptor.
- C. It extracts the first connection request on the queue of pending connections for the socket descriptor.
- D. Assign a socket address to a socket descriptor.

- ☐ Answer A.
- ☐ Answer B.
- ☐ Answer C.
- ✓ ☒ Answer D.

▶ [View Feedback](#)

Question 21

1 / 1 point

Calling `listen(sockfd, backlog)` lets the kernel know that we intend to use `sockfd` as a server-type socket. What is the role of `backlog`?

- A. This is the maximum number of connections that `sockfd` can receive before automatically closing down (or `0` for no limit).
- B. This is a value that indicates whether the kernel should log connections using the standard logging facilities (`/var/log` or `journald`).
- C. This is the size of the buffer for the TCP stream; clients that try to write when that buffer is full will be blocked until the buffer has room.
- D. This is the maximum length of the queue of pending connections to `sockfd`.

☐ Answer A.

☐ Answer B.

☐ Answer C.

☒ Answer D.

▶ [View Feedback](#)

Question 22

1 / 1 point

In a server, what is the return value of a call to `accept(listenfd, ...)`, and how does one communicate with the new client?

- A. It is a new file descriptor that directly connects to the client. The FD `listenfd` remains open and can accept more connections.
- B. It is a new file descriptor that can be used to further accept more connections. The FD `listenfd` is now directly bound to the new client.
- C. It is either `0` if `accept` succeeded, or `-1` otherwise. The FD `listenfd` is used to directly communicate with the client.

☒ Answer A.

☐ Answer B.

☐ Answer C.

▶ [View Feedback](#)

Question 23

1 / 1 point

After a successful call to connect, a process-to-process connection is created. Usually, how is the TCP port on the *client* side chosen?

- A. It is based on the MAC address of the client.
- B. An *ephemeral* port is picked by the kernel.
- C. The server returns the port that the client should use.
- D. A call to bind should have been made before to link the client socket to a port.

☐ Answer A.☒ Answer B.☐ Answer C.☐ Answer D.[▶ View Feedback](#)

Question 24

1 / 1 point

In a client, what is the return value of a call to connect(clientfd, ...), and how does one communicate with the server after the call?

- A. It is a new file descriptor that directly connects to the server. The FD clientfd remains open and can be used to make more connections.
- B. It is either 0 if connect succeeded, or -1 otherwise. The FD clientfd is used to directly communicate with the server.
- C. It is a new file descriptor that can be used to make further connections. The FD clientfd is now directly bound to the server.

☐ Answer A.☒ Answer B.☐ Answer C.[▶ View Feedback](#)

Question 25**1 / 1 point**

We use `telnet` as an echo client (sends the user input raw to the server, and prints the server output as it comes). But this wasn't why `telnet` was created; what was it?

- A. It was used to navigate DNS servers.
- B. It was used as an elementary HTTP client.
- C. It was used as a command-line interface to remote hosts.
- D. It was used as an elementary FTP client.

☐ Answer A.☐ Answer B.☒ Answer C.☐ Answer D.[▶ View Feedback](#)**Question 26****1 / 1 point**

What are the most basic functions used to read and write to a socket?

- A. `sockread` and `sockwrite`.
- B. `read` and `write`.
- C. `rio_readn` and `rio_writen`.
- D. `fread` and `fwrite`.

☐ Answer A.☒ Answer B.☐ Answer C.☐ Answer D.[▶ View Feedback](#)

Question 27**1 / 1 point**

Can the HTTP protocol transfer any type of file?

- A. Yes, but the files need to be transcoded to be using ASCII only (using, e.g., base64).
- B. No, it can only transfer HTML files.
- C. No, it can only transfer files that contain only ASCII characters, and transcoding to ASCII is not part of the protocol.
- D. Yes, the files are sent raw over the connection.

☐ Answer A.

☐ Answer B.

☐ Answer C.

☒ Answer D.

▶ [View Feedback](#)

Question 28**1 / 1 point**

What is the distinction between *static* and *dynamic* content in the context of web servers?

- A. Static is just a file on the server, dynamic is the output of an executable run on the server.
- B. Static is a webpage whose location does not change over time, while dynamic content can be at any location.
- C. Static is a webpage that does not change when it is refreshed, while dynamic pages always change.
- D. Static is a webpage without Javascript, dynamic is a webpage with Javascript.

☒ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

Question 29

1 / 1 point

Assume a web server serves websites `lorem.com` and `ipsum.net`. Recall that a GET request only provides the URL suffix, for instance:

```
GET /mypage.html HTTP/1.0
```

How does the server know whether to serve `lorem.com/mypage.html` or `ipsum.net/mypage.html`?

- A. The GET request should indicate the website domain name as an extra header.
- B. The webserver has one FD per website, so it knows which one the client was addressing.
- C. The webserver knows what was the domain name the client used by inspecting the result of the accept function.
- D. The server returns an error and the GET request is resent by the client with the full URL.

- ✓ ☒ Answer A.
- ☐ Answer B.
- ☐ Answer C.
- ☐ Answer D.

▶ [View Feedback](#)

Question 30

1 / 1 point

What is CGI (Common Gateway Interface)?

- A. A webpage that is dynamically generated.
- B. An HTTP error, indicating that the connection gateway is busy.
- C. A standard for transferring info between the browser and a script that is run on the web server.
- D. A raw format for images, with no compression, used in the early years of the Internet.

- ☐ Answer A.
- ☐ Answer B.
- ✓ ☐ Answer C.
- ☐ Answer D.

▶ [View Feedback](#)

Question 31

1 / 1 point

What is the standard way for passing arguments to a webpage with a GET request?

- A. There is no way to do this with GET, one must use a POST request.
- B. The GET request is followed by request headers of type Argument:, one for each argument.
- C. The GET request takes an optional fourth argument which is a list of arguments, separated by ampersands, e.g.: GET / HTTP/1.1 arg1&arg2.
- D. The URI in the GET request is suffixed with a question mark, then some arguments separated by ampersands.

- ☐ Answer A.
- ☐ Answer B.
- ☐ Answer C.
- ✓ ☐ Answer D.

▶ [View Feedback](#)

Question 32

1 / 1 point

What is the difference between a POST request and a GET request?

- A. GET requests cannot be used with HTML <form>s.
- B. POST requests can send unlimited raw data to the server, GET cannot.
- C. POST requests cannot provide arguments in the URL.
- D. GET requests can request dynamic content, POST cannot.
- E. More than one of the above.

☐ Answer A.

✓ ☒ Answer B.

☐ Answer C.

☐ Answer D.

☐ Answer E.

▶ [View Feedback](#)

Question 33

1 / 1 point

How does a CGI script retrieve the arguments provided in a GET request?

- A. By the usual command line argument passing system (argc and argv in C).
- B. By reading the environment variable QUERY_STRING.
- C. By reading the file /var/cgiargs.
- D. By reading a global variable `extern char *cgiargs`.

☐ Answer A.

✓ ☒ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

Question 34

1 / 1 point

How does a CGI script read the data provided by a POST request?

- A. By reading the file whose filename is stored in the environment variable `CGI_DATA`.
- B. By reading the file `/var/cgidata`.
- C. By reading the FD stored in the global variable `extern int cgidata`.
- D. By reading the standard input.

☐ Answer A.

☐ Answer B.

☐ Answer C.

☒ Answer D.

[▶ View Feedback](#)

Concurrent Programming

Question 35

1 / 1 point

Why is it so hard to picture the interactions within n concurrent executions?

- A. Because the number of possible interactions is exponential in n .
- B. Because it's hard to remember which function calls are `libc`, and which are system calls.
- C. Because it's hard to keep track of the tree of forks (children processes).
- D. Because (single-core) processors only *emulate* concurrency.

☒ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

[▶ View Feedback](#)

Question 36**1 / 1 point**

What are examples of problems that arise in concurrent programming?

- A. Deadlocks: when a process waits for an event that will never happen.
- B. Race conditions: the outcome depends on the order in which processes are scheduled.
- C. Starvation: when a process is never scheduled.
- D. Two of A, B, C.
- E. All of A, B, C.

☐ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

☒ Answer E.

▶ [View Feedback](#)

Question 37**1 / 1 point**

Iterative servers can only serve one client at a time. If a second client tries to connect and make small interactions with the server while it is not in accept mode, when will the second client block?

- A. When calling socket.
- B. When calling connect.
- C. When calling write.
- D. When calling read.

☐ Answer A.

☐ Answer B.

☐ Answer C.

☒ Answer D.

 [View Feedback](#)

Question 38

1 / 1 point

In a process-based concurrent server, each client is assigned a child process (client process). Who is in charge of scheduling these processes? In other terms, what part of the system chooses which process is executed at any given time?

- A. Client processes are normal processes and are scheduled by the kernel with no special consideration.
- B. The CPU is in charge of scheduling client processes as soon as data from them is received.
- C. Client processes are scheduled by the kernel in the order in which the server receives data from them.
- D. The server chooses which client process is to be scheduled next, through system calls.

✓ ☒ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

 [View Feedback](#)

Question 39

1 / 1 point

In a process-based concurrent server, how are processes for each client created?

- A. Using `sockclient`.
- B. Using threads.
- C. Using `execve`.
- D. Using `fork`.

☐ Answer A.

☐ Answer B.

☐ Answer C.☒ Answer D.[▶ View Feedback](#)**Question 40****1 / 1 point**

In a process-based concurrent server, after creating a child process for a new client, the server should close the socket descriptor (file descriptor) associated with that client. Why?

- A. This is optional, and not doing it would not lead to any problem.
- B. If the server were not to do it, when the child process closes that socket the kernel will think that a process still needs that file descriptor, and won't actually close it.
- C. If the server were not to do it, the FD would appear in both the client and the server, but only one of the two processes can read from it.
- D. Two of A, B, C.
- E. All of A, B, C.

☐ Answer A.☒ Answer B.☐ Answer C.☐ Answer D.☐ Answer E.[▶ View Feedback](#)

Attempt Score: 40 / 40**Overall Grade (highest attempt): 40 / 40**

Done