

Timothy Holmes (username: THOLME15)



## Attempt 1

Written: Oct 28, 2020 10:22 PM - Oct 28, 2020 10:31 PM

## Submission View

Released: Sep 7, 2020 10:15 PM

### Question 1

0 / 1 point

In an event-based server, how many different processes/threads are run at any given time?

- A. 1.
- B. 2.
- C. As many as the event poll size.
- D. As many as there are clients.

➡ ☐ Answer A.

☐ Answer B.

☐ Answer C.

✖ ☐ Answer D.

▶ [View Feedback](#)

### Question 2

0 / 1 point

In an event-based server, the server probes all the connections to check if some data is available. How does it do that?

- A. Using a signal handler that is triggered when new data arrives on an open file descriptor.
- B. Using a virtual file descriptor that aggregates all the file descriptors that the process wants to read from.
- C. Using some polling system calls that can wait for an event to arrive on a set of file descriptors.
- D. Using non-blocking sockets, that is, sockets created using the option `O_NONBLOCK`. These sockets return from read immediately if no data is available.

✗ ☐ Answer A.

☐ Answer B.

➡ ☒ Answer C.

☐ Answer D.

▶ [View Feedback](#)

### Question 3

1 / 1 point

In POSIX threads, how is a thread identified?

- A. By both a *thread* ID and the ID of the process in which it lives.
- B. By a *thread* ID that is unique system-wide.
- C. Threads do not have identifiers.

✓ ☒ Answer A.

☐ Answer B.

☐ Answer C.

▶ [View Feedback](#)

### Question 4

0 / 1 point

At any given time, a process state can be described by the current value of the CPU registers, its memory (or more precisely, the virtual memory structures), and the process-specific info in the kernel (e.g., descriptor table). In a process that has two threads, what information is common to the threads?

- A. The memory and the kernel info.
- B. The memory only.
- C. The kernel context only.
- D. All of the info.

➡ ☐ Answer A.

☐ Answer B.

☐ Answer C.

✖ ☐ Answer D.

▶ [View Feedback](#)

### Question 5

0 / 1 point

Threads are created using `pthread_create()`. What is the function used to reap a thread?

- A. `pthread_wait()`.
- B. `pthread_reap()`.
- C. `pthread_join()`.
- D. Threads are not reaped.

- ☐ Answer A.
- ☒ Answer B.
- ☒ Answer C.
- ☐ Answer D.

▶ [View Feedback](#)

## Question 6

0 / 1 point

What are common points between threads and processes?

- A. The costs of creating threads and processes are the same.
- B. The ways memory from the callee thread/process is shared with new threads/processes (created using `pthread_create` and `fork`) are the same.
- C. They have separate control flow.
- D. They run concurrently and are context switched.
- E. Two of A, B, C, D.
- F. Three of A, B, C, D.
- G. All of A, B, C, D.

- ☐ Answer A.
- ☐ Answer B.
- ☐ Answer C.
- ☐ Answer D.
- ☒ Answer E.
- ☐ Answer F.
- ☒ Answer G.

▶ [View Feedback](#)

### Question 7

1 / 1 point

When creating a new thread, the fourth argument of `pthread_create()` is passed as an argument to the thread main function. What is the type of that argument?

- A. `void`.
- B. `int`.
- C. `unsigned long`.
- D. `void *`.

☐ Answer A.

☐ Answer B.

☐ Answer C.

✓ ☒ Answer D.

▶ [View Feedback](#)

### Question 8

0 / 1 point

What is detached mode?

- A. A running mode for processes in which they can't reap the threads they spawn.
- B. A running mode for threads in which they don't need (and can't) be reaped.
- C. A running mode for threads in which they can't reap other threads.
- D. A running mode for processes in which they don't need (and can't) be reaped.

☐ Answer A.

➡ ☒ Answer B.

✗ ☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

### Question 9

0 / 1 point

Why is it hard, in practice, to identify which variables are shared between threads?

- A. Because any thread can access the whole memory space of the process, including the stacks of the other threads.
- B. Because we don't have control over the scheduling of threads, as this is the kernel's job.
- C. Because threads don't share the same memory space.
- D. More than one of A, B, C.

➡ ☐ Answer A.

☐ Answer B.

☐ Answer C.

✗ ☐ Answer D.

▶ [View Feedback](#)

### Question 10

0 / 1 point

Assume we want to pass a few parameters when creating a handful of threads. We bundle these arguments into a **struct** that contains, for instance, a thread number. Instead of passing the address of the structure directly in the fourth argument of `pthread_open()`, we first make a copy, using `malloc`, and pass this copy, that the new thread frees as soon as it has processed its arguments. Why didn't we pass the original structure in the first place?

- A. Because the new thread would access a value that is located on the main thread's stack, resulting in a memory fault.
- B. Because the original structure could be manipulated by the main thread before it is read by the new thread.
- C. Because it would be less efficient.
- D. More than one of A, B, C.

✗ ☐ Answer A.

➡ ☒ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

## Question 11

1 / 1 point



What is the difference between a global variable and a local static variable?

- A. The number of instances of each variable in threads: global variables appear exist once for all the threads, local static variables have one instance per thread.
- B. Where the variable is allocated: global variables are allocated in the data or bss segment, while local static variables are allocated on the stack.
- C. The scope of the variable: global variables can be used everywhere, while local static variables are only accessible within the scope of their definition.
- D. More than one of A, B, C.

☐ Answer A.

☐ Answer B.

☒ Answer C.

☐ Answer D.

▶ [View Feedback](#)

## Question 12

0 / 1 point

What is the role of the **volatile** keyword?

- A. It is used in function definitions. It indicates to the compiler that the function will have a side effect.
- B. It is used in variable definition. It forces all references to the variable to be made from memory, rather than storing it in a register.
- C. It is used in variable definition. It asks the compiler to deallocate the variable as soon as it detects it is not referenced anymore, to free up space.
- D. It is used in function definitions. It asks the compiler to never *inline* the function (a function is *inlined* if calls to it are replaced by the actual body of the function, in the resulting assembly).



✗ ☐ Answer A.

➡ ☐ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

### Question 13

0 / 1 point

Why isn't incrementing a value in memory thread-safe?

- A. Because it consists of three phases: fetch the value in a register, increment the register, then write it back to memory. A thread can be interrupted in the middle, and, when resumed, have an outdated value in the register.
- B. Because accessing a value in memory can cause a page fault exception, and during the loading of the page in memory, another thread could be executed.
- C. Because accessing a value in memory can cause a page fault exception, and block all the other threads until the page is loaded in memory, resulting in poor performances.

➡ ☐ Answer A.

☐ Answer B.

✗ ☐ Answer C.

▶ [View Feedback](#)

### Question 14

0 / 1 point

Semaphores are used in particular for mutual exclusion, but what is a semaphore, in terms of type? In other words, what is `sem_t`?

- A. A structure.
- B. A `char`.
- C. An integer.
- D. A pointer.

✗ ☐ Answer A.

☐ Answer B.

➡ ☒ Answer C.

☐ Answer D.

▶ [View Feedback](#)

## Question 15

1 / 1 point

The two operations on a semaphore are traditionally called P and V. What are their roles?

- A. P(s) suspends execution until s is nonzero, then decrements it. V(s) increments s.
- B. V(s) tests if s is zero, and increments it if it is, decrements it otherwise. P(s) zeroes out s.
- C. V(s) suspends execution until s is nonzero, then decrements it. P(s) increments s.
- D. P(s) tests if s is zero, and increments it if it is, decrements it otherwise. V(s) zeroes out s.

✓ ☒ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

### Question 16

0 / 1 point

Why are P and V implemented using system calls, rather than just normal memory manipulations?

A. Because they reference a variable shared across threads.

B. To ensure atomicity of their operations.

C. To avoid page faults.

D. To make them faster.

☐ Answer A.

➡ ☐ Answer B.

✗ ☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

### Question 17

1 / 1 point

In the course, we use P and V, and these are implemented using (actually, just wrappers around) some functions implemented in the libc. Which ones?

- A. P is a wrapper around `sem_wait`, and V is a wrapper around `sem_post`.
- B. V is a wrapper around `sem_post`, and P is a wrapper around `sem_init`.
- C. P is a wrapper around `sem_post`, and V is a wrapper around `sem_wait`.
- D. V is a wrapper around `sem_init`, and P is a wrapper around `sem_post`.

✓ ☒ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

## Question 18

1 / 1 point

Why is **volatile** (on global variables) considered harmful in multithreaded applications?

- A. Because it is useless in practice, since library calls will prevent global variables to be put in registers.
- B. Because it prevents a lot of optimizations by the compiler.
- C. Because it provides a false sense of atomicity.
- D. Two of A, B, C.
- E. All of A, B, C.

☐ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

✓ ☐ Answer E.

▶ [View Feedback](#)

### Question 19

1 / 1 point

In the theory of semaphores, what is a mutex?

- A. A semaphore taking only values 0 and 1 used for mutual exclusion (P locks the mutex, V releases it).
- B. Mutexes are different from semaphores, and semaphores cannot implement mutexes.
- C. A design pattern that uses semaphores to avoid concurrent execution of critical portions of code.
- D. A type of semaphore implemented using a single bit.

✓ ☐ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

▶ [View Feedback](#)

### Question 20

0 / 1 point

In the producer-consumer problem, a pool of producers puts data in a fixed-size buffer, and a pool of consumer reads and removes data from the buffer. What is the problem that semaphores solves?

- A. How to wait for data to arrive.
- B. How to schedule producers and consumers.
- C. Concurrent access to the buffer, so that it is in a consistent state at all time.
- D. How to wait for free space to be available in the buffer.
- E. Two of A, B, C, D.
- F. Three of A, B, C, D.
- G. All of A, B, C, D.

☐ Answer A.

☐ Answer B.

☐ Answer C.

☐ Answer D.

☐ Answer E.

☒ Answer F.

☐ Answer G.

▶ [View Feedback](#)

## Question 21

0 / 1 point



What are the semaphores used solving the producer-consumer problem?

- A. Just one mutex for accessing the shared buffer.
- B. Two counting semaphores to signal production and consumption.
- C. One mutex for modifying the buffer, and one counting semaphore to signal both production and consumption.
- D. One mutex for modifying the buffer, and two counting semaphores to signal production and consumption.

☐ Answer A.

☐ Answer B.

☒ Answer C.

☐ Answer D.

▶ [View Feedback](#)

## Question 22

1 / 1 point

In the readers-writers problem, some reader threads want to read a shared object, and some writer threads want to write to it. This could be solved using a single mutex, ensuring that no two threads have access to the object at the same time. Why is this very much suboptimal?

- A. Because multiple writers should be allowed to write to the object at the same time.
- B. Because one writer and one reader should be allowed to access the object at the same time.
- C. Because multiple readers should be allowed to read the object at the same time.

☐ Answer A.

☐ Answer B.

✓ ☐ Answer C.

▶ [View Feedback](#)

---

**Attempt Score:**8 / 22

**Overall Grade (highest attempt):**8 / 22

Done