

# Concurrent Programming:

## *Concurrent Servers:*

### *3. Thread based*

# Approaches for Writing Concurrent Servers

Allow server to handle multiple clients concurrently

## 1. Process-based

- Kernel automatically interleaves multiple logical flows
- Each flow has its own private address space

## 2. Event-based

- Programmer manually interleaves multiple logical flows
- All flows share the same address space
- Uses technique called *I/O multiplexing*.

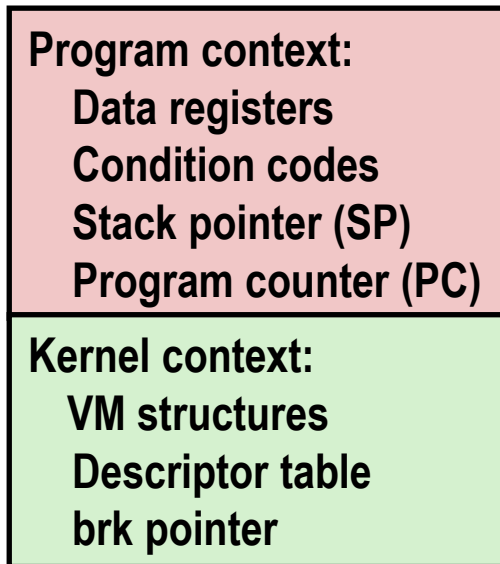
## 3. Thread-based

- Kernel automatically interleaves multiple logical flows
- Each flow shares the same address space
- Hybrid of of process-based and event-based.


# Traditional View of a Process

- **Process = process context + code, data, and stack**

## Process context

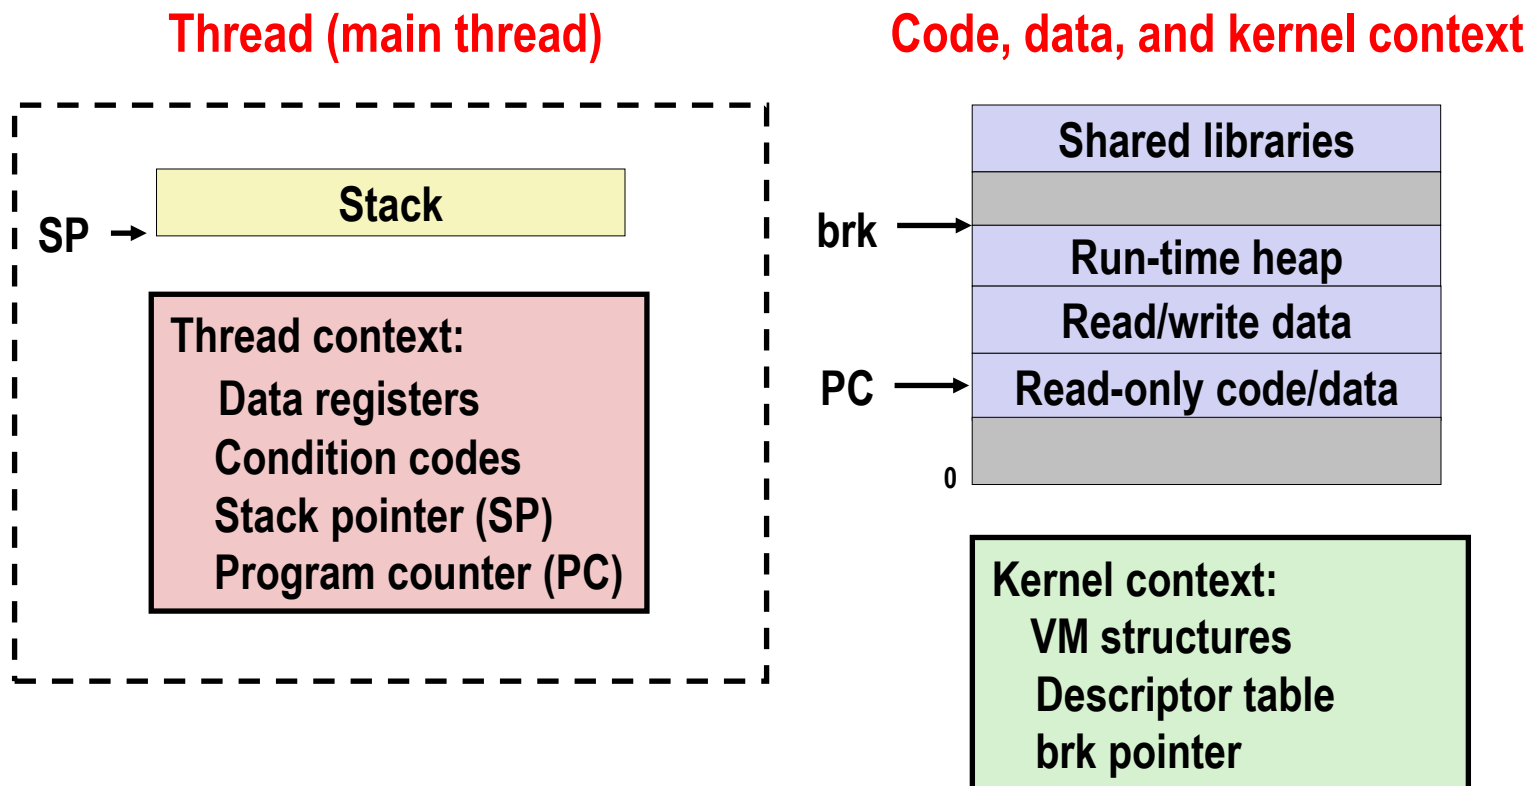


## Code, data, and stack



# Alternate View of a Process

- Process = thread + code, data, and kernel context



# A Process With Multiple Threads

- **Multiple threads can be associated with a process**
  - Each thread shares the same code, data, and kernel context
  - Each thread has its own logical control flow
  - Each thread has its own stack for local variables
    - but not protected from other threads
  - Each thread has its own thread id (TID), unique within process

**Thread 1 (main thread)**

**Thread 2 (peer thread)**


**Shared code and data**

**stack 1**

**stack 2**

Thread 1 context:  
Data registers  
Condition codes  
SP1  
PC1

Thread 2 context:  
Data registers  
Condition codes  
SP2  
PC2



Kernel context:  
VM structures  
Descriptor table  
brk pointer


- ***Lower cost of context switch!***

# Logical View of Threads


## ■ Threads associated with process form a *pool* of peers

- Unlike processes which form a tree hierarchy
- Possibility of mix-and-matching, but gets complex

Threads associated with process foo



Process hierarchy



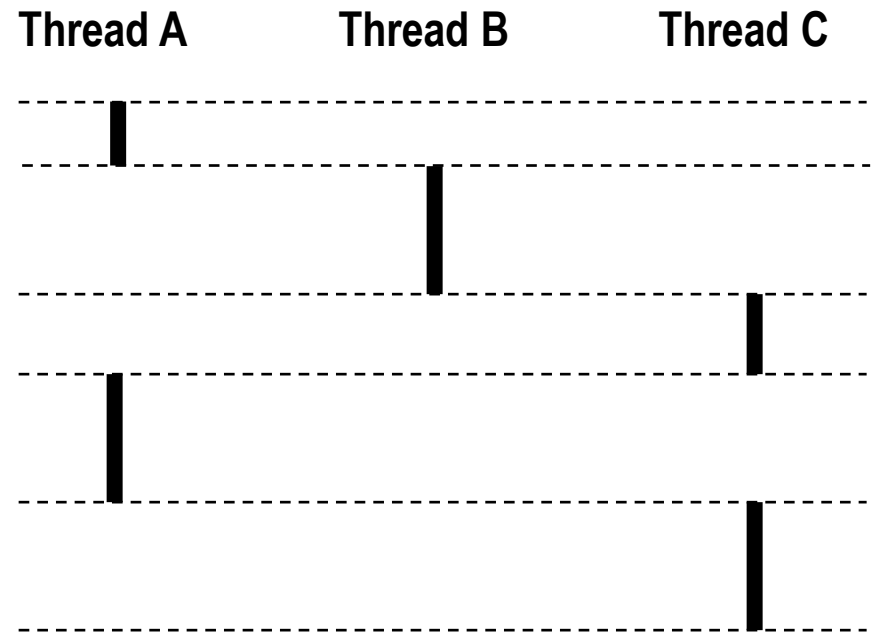
# Concurrent Threads

- Two threads are *concurrent* if their flows overlap in time
- Otherwise, they are sequential

- **Examples:**

- Concurrent: A & B, A&C
- Sequential: B & C

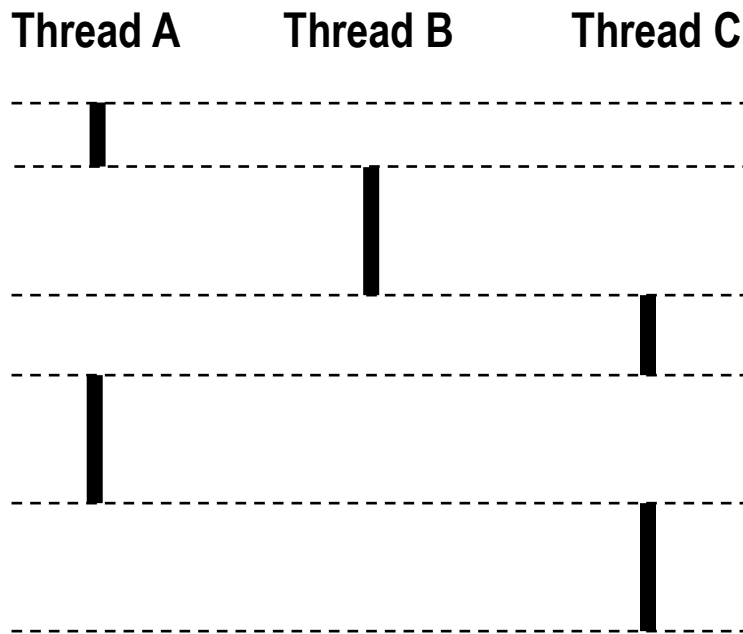
Time



# Concurrent Thread Execution


## ■ Single Core Processor

- Simulate parallelism by time slicing



## ■ Multi-Core Processor

- Can have true parallelism





# Threads vs. Processes

## ■ How threads and processes are similar

- Each has its own logical control flow
- Each can run concurrently with others (possibly on different cores)
- Each is context switched

## ■ How threads and processes are different

- Threads share all code and data (except local stacks)
  - Processes (typically) do not
- Threads are somewhat less expensive than processes
  - Process control (creating and reaping) twice as expensive as thread control
  - Linux numbers:
    - ~20K cycles to create and reap a process
    - ~10K cycles (or less) to create and reap a thread

# Posix Threads (Pthreads) Interface

- ***Pthreads*: Standard thread interface (~60 functions)**
  - Creating and reaping threads
    - `pthread_create(3)`
    - `pthread_join(3)`
  - Determining your thread ID
    - `pthread_self(3)`
  - Terminating threads
    - `pthread_cancel(3)`
    - `pthread_exit(3)`
    - `exit(3)` [terminates all threads]
  - Synchronizing access to shared variables
    - `pthread_mutex_init(3)`
    - `pthread_mutex_[un]lock(3)`
- **libc implemented (`-lpthread`), syscall is `clone(2)`**
  - See <https://nullprogram.com/blog/2015/05/15/> for raw calls

# The Pthreads "hello, world" Program

```
/*  
 * hello.c - Pthreads "hello, world" program  
 */  
#include "csapp.h"  
void *thread(void *vargp);  
  
int main()  
{  
    pthread_t tid;  
    Pthread_create(&tid, NULL, thread, NULL);  
    Pthread_join(tid, NULL);  
    exit(0);  
}  
hello.c
```

```
void *thread(void *vargp) /* thread routine */  
{  
    printf("Hello, world!\n");  
    return NULL;  
}  
hello.c
```

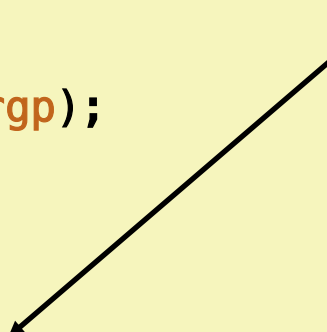
# The Pthreads "hello, world" Program

```
/*
 * hello.c - Pthreads "hello, world" program
 */
#include "csapp.h"
void *thread(void *vargp);

int main()
{
    pthread_t tid;
    Pthread_create(&tid, NULL, thread, NULL);
    Pthread_join(tid, NULL);
    exit(0);
}
```

Thread ID

hello.c



```
void *thread(void *vargp) /* thread routine */
{
    printf("Hello, world!\n");
    return NULL;
}
```

hello.c

# The Pthreads "hello, world" Program

```
/*
 * hello.c - Pthreads "hello, world" program
 */
#include "csapp.h"
void *thread(void *vargp);

int main()
{
    pthread_t tid;
    Pthread_create(&tid, NULL, thread, NULL);
    Pthread_join(tid, NULL);
    exit(0);
}
```

hello.c

Thread ID

*Thread attributes  
(usually NULL,  
see pthread\_attr\_init(3))*

```
void *thread(void *vargp) /* thread routine */
{
    printf("Hello, world!\n");
    return NULL;
}
```

hello.c

# The Pthreads "hello, world" Program

```
/*
 * hello.c - Pthreads "hello, world" program
 */
#include "csapp.h"
void *thread(void *vargp);

int main()
{
    pthread_t tid;
    Pthread_create(&tid, NULL, thread, NULL);
    Pthread_join(tid, NULL);
    exit(0);
}
```

hello.c

Thread ID

Thread attributes  
(usually NULL,  
see pthread\_attr\_init(3))

Thread routine

```
void *thread(void *vargp) /* thread routine */
{
    printf("Hello, world!\n");
    return NULL;
}
```

hello.c

# The Pthreads "hello, world" Program

```
/*  
 * hello.c - Pthreads "hello, world" program  
 */  
#include "csapp.h"  
void *thread(void *vargp);  
  
int main()  
{  
    pthread_t tid;  
    Pthread_create(&tid, NULL, thread, NULL);  
    Pthread_join(tid, NULL);  
    exit(0);  
}
```

hello.c

Thread ID

*Thread attributes  
(usually NULL,  
see pthread\_attr\_init(3))*

Thread routine

*Thread arguments  
(void \*p)*

```
void *thread(void *vargp) /* thread routine */  
{  
    printf("Hello, world!\n");  
    return NULL;  
}
```

hello.c

# The Pthreads "hello, world" Program

```
/*
 * hello.c - Pthreads "hello, world" program
 */
#include "csapp.h"
void *thread(void *vargp);

int main()
{
    pthread_t tid;
    Pthread_create(&tid, NULL, thread, NULL);
    Pthread_join(tid, NULL);
    exit(0);
}
```

hello.c

Thread ID

Thread attributes  
(usually NULL,  
see pthread\_attr\_init(3))

Thread routine

Thread arguments  
(void \*p)

Return value  
(void \*\*p)

```
void *thread(void *vargp) /* thread routine */
{
    printf("Hello, world!\n");
    return NULL;
}
```

hello.c




# Execution of Threaded “hello, world”


Main thread




# Execution of Threaded “hello, world”




# Execution of Threaded “hello, world”




# Execution of Threaded “hello, world”



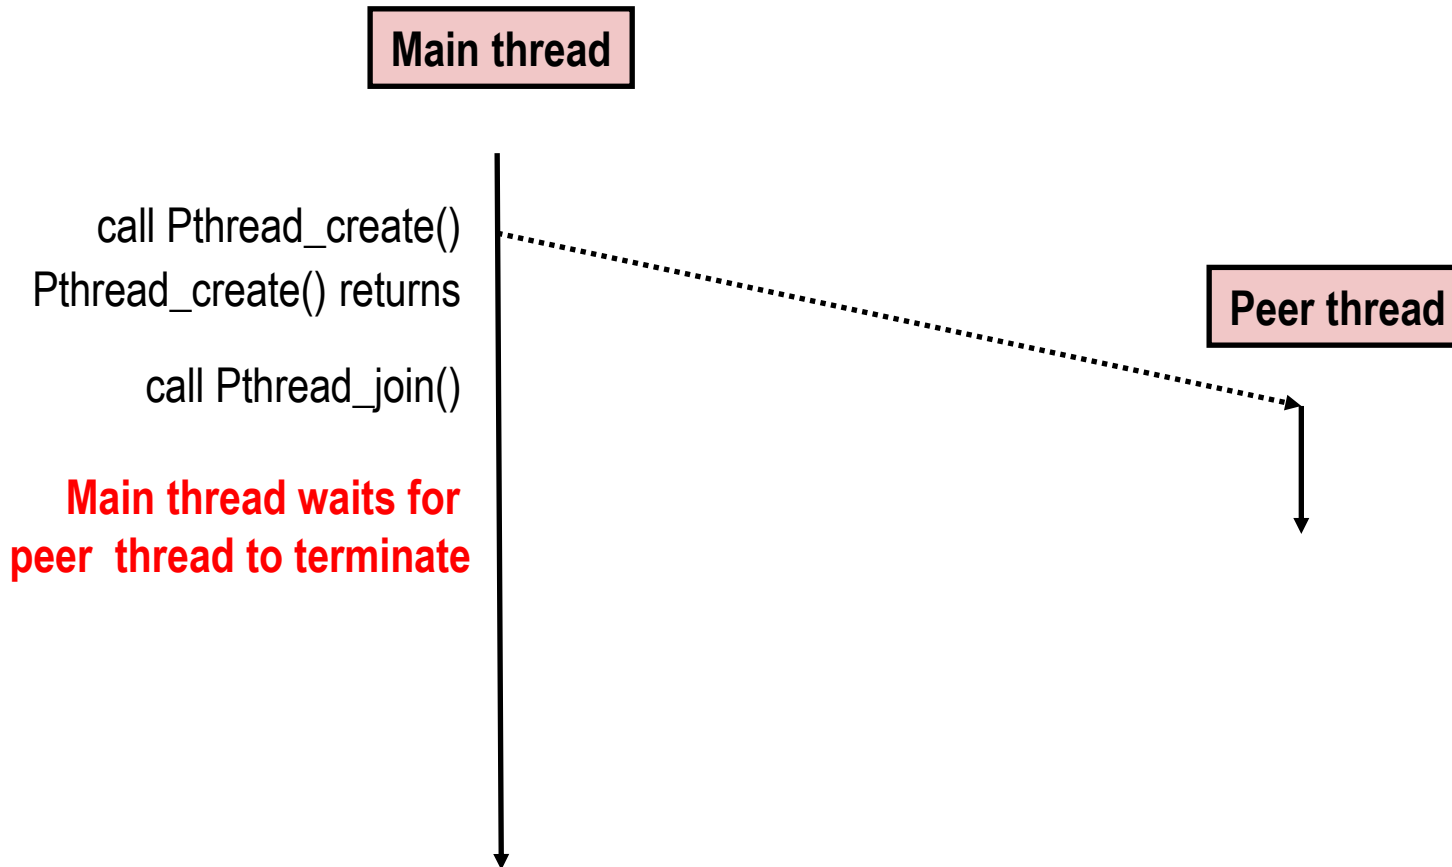
# Execution of Threaded “hello, world”



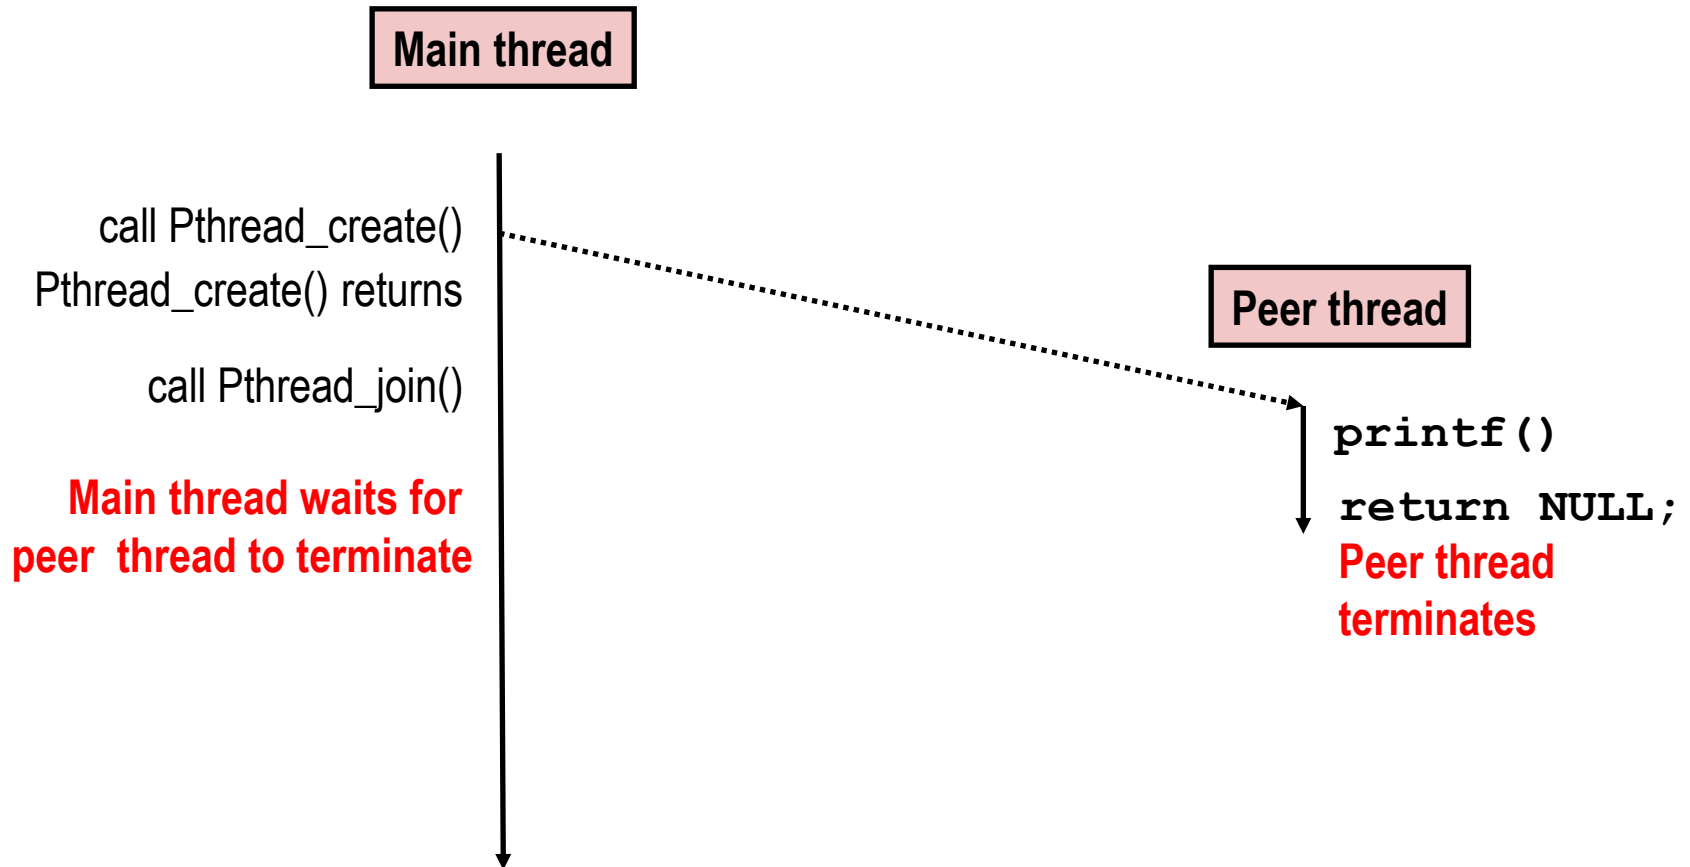
# Execution of Threaded “hello, world”



# Execution of Threaded “hello, world”

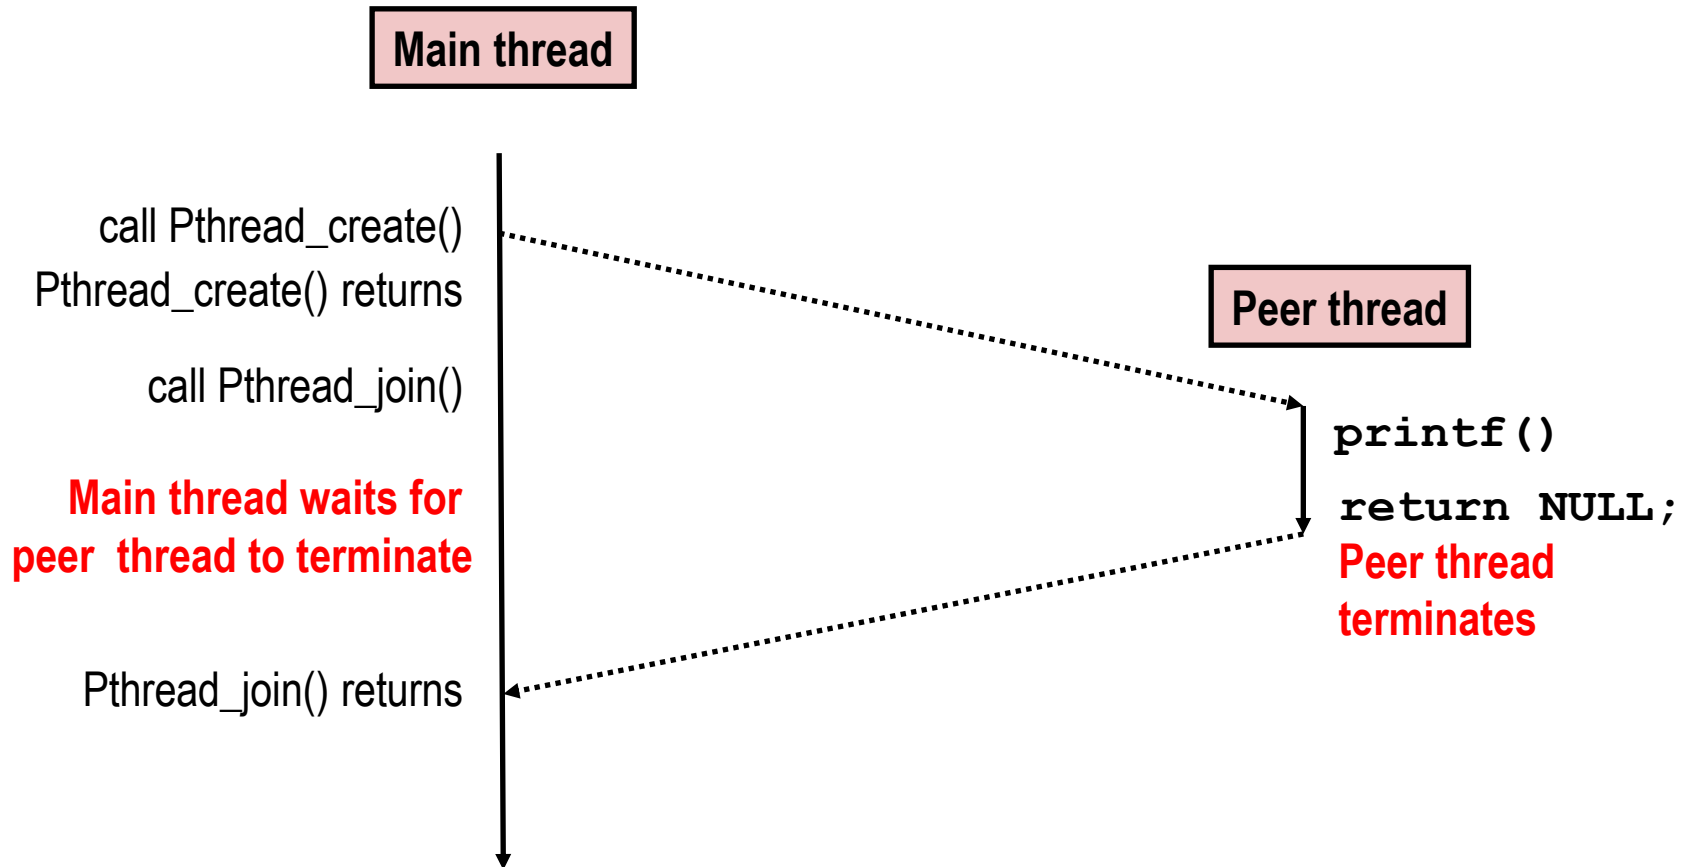


# Execution of Threaded “hello, world”

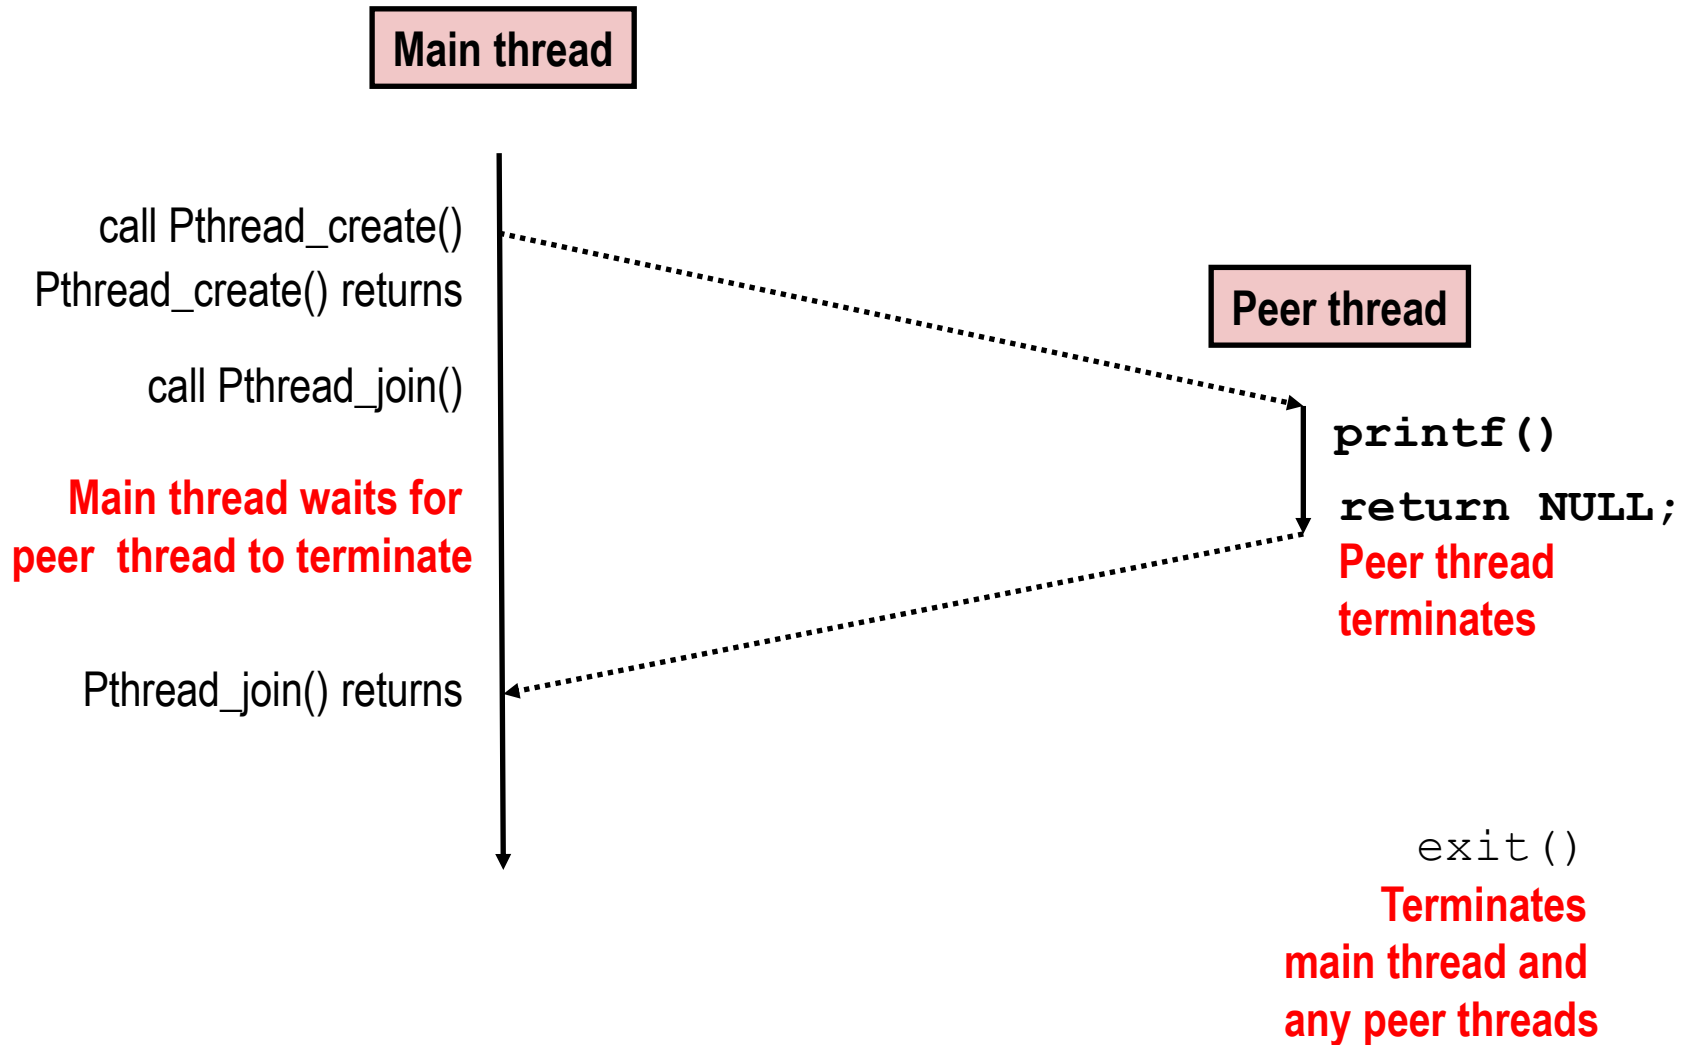




# Execution of Threaded “hello, world”



# Execution of Threaded “hello, world”



# Thread-Based Concurrent Echo Server

```
int main(int argc, char **argv)
{
    int listenfd, *connfdp;
    socklen_t clientlen;
    struct sockaddr_storage clientaddr;
    pthread_t tid;

    listenfd = Open_listenfd(argv[1]);
    while (1) {
        clientlen = sizeof(struct sockaddr_storage);
        connfdp = Malloc(sizeof(int));
        *connfdp = Accept(listenfd,
                          (SA *) &clientaddr, &clientlen);
        Pthread_create(&tid, NULL, thread, connfdp);
    }
}
```

echoserver.c

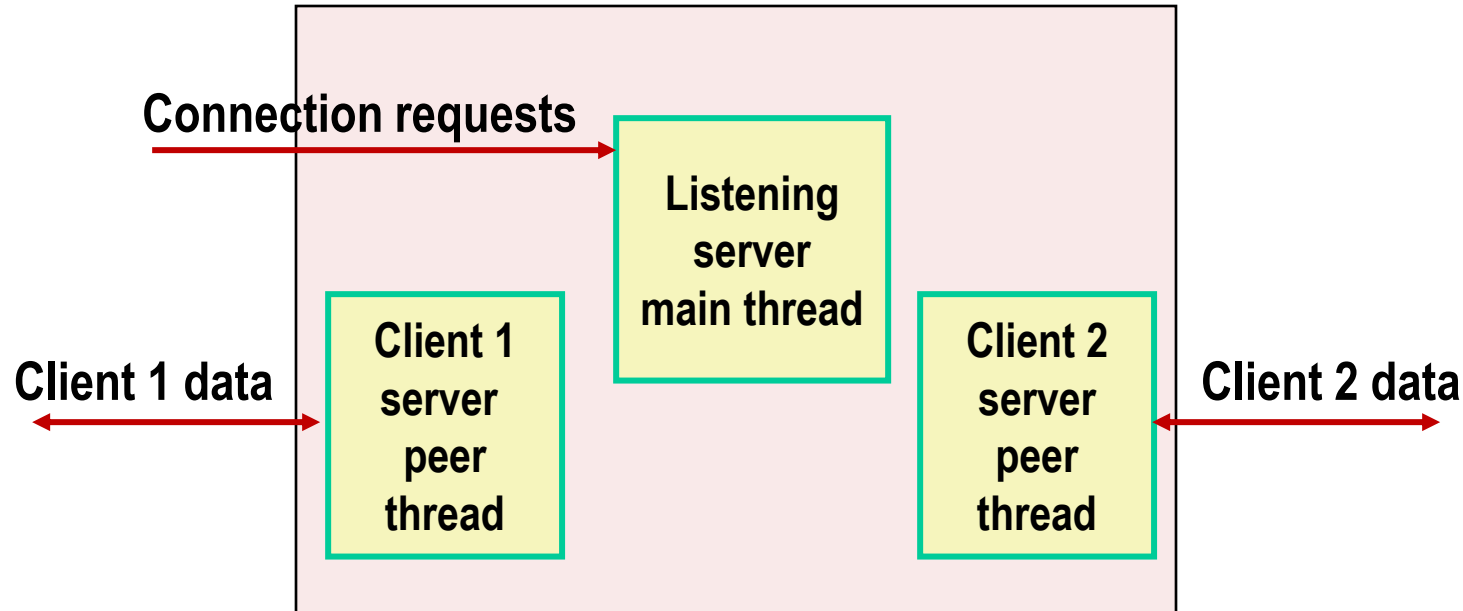
- malloc of connected descriptor necessary to avoid deadly race (later)
- Should NOT close connfd!

# Thread-Based Concurrent Server (cont)

```
/* Thread routine */
void *thread(void *vargp)
{
    int connfd = *((int *)vargp);
    Pthread_detach(pthread_self());
    Free(vargp);
    echo(connfd);
    Close(connfd);
    return NULL;
}
echoserv.c
```

- Run thread in “detached” mode.
  - Can’t be made joinable again
  - Reaped automatically (by kernel) when it terminates
  - Either `pthread_detach` or `pthread_join` should be called for each thread, but never both!
- Must free storage allocated to hold `connfd`.
- Must close `connfd` (important, shared FD table!)

# Thread-based Server Execution Model



- Each client handled by individual peer thread
- Threads share all process state in kernel, except TID
- Each thread has a separate stack for local variables, and own registers

# Issues With Thread-Based Servers

## ■ Must join or detach

- At any point in time, a thread is either *joinable* or *detached*
- *Joinable* thread can be reaped and killed by other threads
  - must be reaped (with `pthread_join`) to free memory resources
- *Detached* thread cannot be reaped or killed by other threads
  - resources are automatically reaped on termination
- Default state is joinable (useful for splitting task: “fork-join model”)
  - use `pthread_detach(pthread_self())` to make detached
  - Or create in detach mode with `pthread_attr_setdetachstate`

# Issues With Thread-Based Servers

## ■ Must join or detach

- At any point in time, a thread is either *joinable* or *detached*
- *Joinable* thread can be reaped and killed by other threads
  - must be reaped (with `pthread_join`) to free memory resources
- *Detached* thread cannot be reaped or killed by other threads
  - resources are automatically reaped on termination
- Default state is joinable (useful for splitting task: “fork-join model”)
  - use `pthread_detach(pthread_self())` to make detached
  - Or create in detach mode with `pthread_attr_setdetachstate`

## ■ Must be careful to avoid unintended sharing

- For example, passing pointer to main thread's stack
  - `Pthread_create(&tid, NULL, thread, (void *)&connfd);`

# Issues With Thread-Based Servers

## ■ Must join or detach

- At any point in time, a thread is either *joinable* or *detached*
- *Joinable* thread can be reaped and killed by other threads
  - must be reaped (with `pthread_join`) to free memory resources
- *Detached* thread cannot be reaped or killed by other threads
  - resources are automatically reaped on termination
- Default state is joinable (useful for splitting task: “fork-join model”)
  - use `pthread_detach(pthread_self())` to make detached
  - Or create in detach mode with `pthread_attr_setdetachstate`

## ■ Must be careful to avoid unintended sharing

- For example, passing pointer to main thread's stack
  - `Pthread_create(&tid, NULL, thread, (void *)&connfd);`
  - Assumes that address dereferenced in peer thread before changed!



# Issues With Thread-Based Servers

## ■ Must join or detach

- At any point in time, a thread is either *joinable* or *detached*
- *Joinable* thread can be reaped and killed by other threads
  - must be reaped (with `pthread_join`) to free memory resources
- *Detached* thread cannot be reaped or killed by other threads
  - resources are automatically reaped on termination
- Default state is joinable (useful for splitting task: “fork-join model”)
  - use `pthread_detach(pthread_self())` to make detached
  - Or create in detach mode with `pthread_attr_setdetachstate`

## ■ Must be careful to avoid unintended sharing

- For example, passing pointer to main thread's stack
  - `Pthread_create(&tid, NULL, thread, (void *)&connfd);`
  - Assumes that address dereferenced in peer thread before changed!

## ■ All functions called by a thread must be *thread-safe*

- Much more on that later

# Pros and Cons of Thread-Based Designs

- **+ Easy to share data structures between threads**
  - e.g., logging information, file cache
- **+ Threads are more efficient than processes**
- **– Unintentional sharing can introduce subtle and hard-to-reproduce errors!**
  - The ease with which data can be shared is both the greatest strength and the greatest weakness of threads
  - Hard to know which data shared & which private
  - Hard to detect by testing
    - Probability of bad race outcome very low
    - But nonzero!
  - Future lectures

# Summary: Approaches to Concurrency

## ■ Process-based

- Hard to share resources/Easy to avoid unintended sharing
- High overhead in adding and removing clients

## ■ Event-based

- Tedious and low level
- Total control over scheduling
- Very low overhead
- Cannot create as fine grained a level of concurrency
- Does not make use of multi-core

## ■ Thread-based

- Easy to share resources... Perhaps too easy!
- Medium overhead
- Not much control over scheduling policies
- Difficult to debug
  - Event orderings not repeatable