

Concurrent Programming:

Synchronizing threads:

2. Semaphores

Semaphores

- ***Semaphore***: non-negative *global (shared)* integer synchronization variable. Manipulated by atomic operations *P* and *V*.

Semaphores

- ***Semaphore***: non-negative *global (shared)* integer synchronization variable. Manipulated by atomic operations *P* and *V*.
- **P(s)**:
 - If *s* is nonzero, then decrement *s* by 1 and return immediately.
 - Test and decrement operations occur atomically (indivisibly)
 - If *s* is zero, then suspend thread until *s* becomes nonzero and the thread is restarted by a *V* operation.
 - After restarting, the *P* operation decrements *s* and returns control to the caller.

Semaphores

- ***Semaphore***: non-negative *global (shared)* integer synchronization variable. Manipulated by atomic operations *P* and *V*.
- **P(s)**:
 - If *s* is nonzero, then decrement *s* by 1 and return immediately.
 - Test and decrement operations occur atomically (indivisibly)
 - If *s* is zero, then suspend thread until *s* becomes nonzero and the thread is restarted by a *V* operation.
 - After restarting, the *P* operation decrements *s* and returns control to the caller.
- **V(s)**:
 - Increment *s* by 1.
 - Increment operation occurs atomically
 - If there are any threads blocked in a *P* operation waiting for *s* to become nonzero, then restart exactly one of those threads, which then completes its *P* operation by decrementing *s*.

Semaphores

- ***Semaphore***: non-negative *global (shared)* integer synchronization variable. Manipulated by atomic operations *P* and *V*.
- **P(s)**:
 - If *s* is nonzero, then decrement *s* by 1 and return immediately.
 - Test and decrement operations occur atomically (indivisibly)
 - If *s* is zero, then suspend thread until *s* becomes nonzero and the thread is restarted by a *V* operation.
 - After restarting, the *P* operation decrements *s* and returns control to the caller.
- **V(s)**:
 - Increment *s* by 1.
 - Increment operation occurs atomically
 - If there are any threads blocked in a *P* operation waiting for *s* to become non-zero, then restart exactly one of those threads, which then completes its *P* operation by decrementing *s*.
- **Semaphore invariant: ($s \geq 0$)**

Semaphores

- **Semaphore:** non-negative *global (shared)* integer synchronization variable. Manipulated by atomic operations *P* and *V*.
- **P(s):**
 - If *s* is nonzero, then decrement *s* by 1 and return immediately.
 - Test and decrement operations occur atomically (indivisibly)
 - If *s* is zero, then suspend thread until *s* becomes nonzero and the thread is restarted by a *V* operation.
 - After restarting, the *P* operation decrements *s* and returns control to the caller.
- **V(s):**
 - Increment *s* by 1.
 - Increment operation occurs atomically
 - If there are any threads blocked in a *P* operation waiting for *s* to become non-zero, then restart exactly one of those threads, which then completes its *P* operation by decrementing *s*.
- **Semaphore invariant:** ($s \geq 0$)
- **Proberen/Verhogen, Procure/Vacate, Acquire/Release, Wait/Signal**

C Semaphore Operations

Pthreads functions, implemented in libc, using CPU-dependent instructions and syscall `futex` (since 2015)

```
#include <semaphore.h>

int sem_init(sem_t *s, 0, unsigned int val); /* s = val */

int sem_wait(sem_t *s); /* P(s) */
int sem_post(sem_t *s); /* V(s) */
```

<https://developers.redhat.com/blog/2015/01/28/recent-improvements-to-concurrent-code-in-glibc/>

CS:APP wrapper functions:

```
#include "csapp.h"

void P(sem_t *s); /* Wrapper function for sem_wait */
void V(sem_t *s); /* Wrapper function for sem_post */
```

badcnt.c: Improper Synchronization

```
/* Global shared variable */
volatile long cnt = 0; /* Counter */

int main(int argc, char **argv)
{
    long niters;
    pthread_t tid1, tid2;

    niters = atoi(argv[1]);
    Pthread_create(&tid1, NULL,
                  thread, &niters);
    Pthread_create(&tid2, NULL,
                  thread, &niters);
    Pthread_join(tid1, NULL);
    Pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%ld\n", cnt);
    else
        printf("OK cnt=%ld\n", cnt);
    exit(0);
}
```

badcnt.c

```
/* Thread routine */
void *thread(void *vargp)
{
    long i, niters =
        *((long *)vargp);

    for (i = 0; i < niters; i++)
        cnt++;

    return NULL;
}
```

How can we fix this using semaphores?

Using Semaphores for Mutual Exclusion

■ Basic idea:

- Associate a unique semaphore *mutex*, initially 1, with each shared variable (or related set of shared variables).
- Surround corresponding critical sections with $P(mutex)$ and $V(mutex)$ operations.

■ Terminology:

- *Binary semaphore*: semaphore whose value is always 0 or 1
- *Mutex*: binary semaphore used for mutual exclusion
 - P operation: “locking” the mutex
 - V operation: “unlocking” or “releasing” the mutex
 - “Holding” a mutex: locked and not yet unlocked.
- *Counting semaphore*: used as a counter for set of available resources.

goodcnt.c: Proper Synchronization

- Define and initialize a mutex for the shared variable `cnt`:

```
volatile long cnt = 0;  /* Counter */
sem_t mutex;           /* Semaphore that protects cnt */

Sem_init(&mutex, 0, 1); /* mutex = 1 */
```

- Surround critical section with *P* and *V*:

```
for (i = 0; i < niters; i++) {
    P(&mutex);
    cnt++;
    V(&mutex);
}
```

goodcnt.c

```
$ ./goodcnt 10000
OK cnt=20000
$ ./goodcnt 10000
OK cnt=20000
$
```

goodcnt.c: Proper Synchronization

- Define and initialize a mutex for the shared variable `cnt`:

```
volatile long cnt = 0;  /* Counter */
sem_t mutex;           /* Semaphore that protects cnt */

Sem_init(&mutex, 0, 1); /* mutex = 1 */
```

- Surround critical section with *P* and *V*:

```
for (i = 0; i < niters; i++) {
    P(&mutex);
    cnt++;
    V(&mutex);
}
```

goodcnt.c

```
$ ./goodcnt 10000
OK cnt=20000
$ ./goodcnt 10000
OK cnt=20000
$
```

Warning: It's orders of magnitude slower than `badcnt.c`.

goodcnt.c: Proper Synchronization

- Define and initialize a mutex for the shared variable cnt:

```
volatile long cnt = 0;  /* Counter */
sem_t mutex;           /* Semaphore that protects cnt */

Sem_init(&mutex, 0, 1); /* mutex = 1 */
```

- Surround critical section with *P* and *V*:

```
for (i = 0; i < niters; i++) {
    P(&mutex);
    cnt++;
    V(&mutex);
}
```


goodcnt.c

```
$ ./goodcnt 10000
OK cnt=20000
$ ./goodcnt 10000
OK cnt=20000
$
```

Warning: It's orders of magnitude slower than badcnt.c. Better put P, V around loop?

Why Mutexes Work

Thread 2



Provide mutually exclusive access to shared variable by surrounding critical section with P and V operations on semaphore s (initially set to 1)

Semaphore invariant creates a **forbidden region** that encloses unsafe region and that cannot be entered by any trajectory.

Thread 1

Initially

$s = 1$

volatile considered useless (or harmful)

```
volatile long cnt = 0; /* Counter */
```

- **volatile** on global variables gives a sense of safety
- ...but *only* gives a hint to compiler to sync variable back to memory after each line of code (costly!)

volatile considered useless (or harmful)

```
volatile long cnt = 0; /* Counter */
```

- **volatile** on global variables gives a sense of safety
- ...but *only* gives a hint to compiler to sync variable back to memory after each line of code (costly!)
- *Race conditions still very much possible!*

volatile considered useless (or harmful)

```
volatile long cnt = 0; /* Counter */
```

- **volatile** on global variables gives a sense of safety
- ...but *only* gives a hint to compiler to sync variable back to memory after each line of code (costly!)
- *Race conditions still very much possible!*
- We do not want **cnt** to end up in a register...
- ...but it won't!

```
{  
    my_library_call (); /* Can't assume won't modify cnt */  
    ...  
}
```